EVIDENCE-BASED HEALTHCARE AI SYSTEM

MedQuAD AI Medical Assistant

Clinical Question Answering, Disease Exploration & Diagnostic Knowledge Retrieval

Presented By:

Rashi Kanwar Rao

Lead Project Developer & Healthcare AI Researcher

Dataset & Foundation:

NIH MedQuAD Benchmark

16,350+ Clinical Q&A Pairs | 5,120+ Conditions

Project Motivation & Clinical Need

⚠ The Healthcare Information Challenge

- Unreliable Public Web Queries: Patients seeking medical advice online are frequently misled by unverified or sensationalized medical advice.
- Clinical Density & Jargon: Authoritative medical literature uses dense terminology (dyspnea, idiopathic, ischemia) that creates cognitive barriers.
- Diagnostic Anxiety: Misinterpreted health information increases patient anxiety and delays timely consultation with licensed professionals.
- Lack of Source Attribution: Most online aggregators do not cite official clinical origins or specific institutional research.

The MedQuAD AI Solution

- NIH-Verified Ground Truth: Grounded strictly in peer-reviewed clinical knowledge from 12 National Institutes of Health (NIH) entities.
- Hybrid Semantic Retrieval: Combines sublinear TF-IDF, n-gram vector matching, and clinical entity boosting for precision answers.
- Patient-Centered NLP Layer: Automatically distills answers into actionable bulleted takeaways and plain-English medical glossaries.
- Responsible Medical Governance: Transparent similarity confidence scores and prominent clinical safety notices.

HEALTHCARE AI & NLP CAPSTONE

MedQuAD Dataset & Institutional Provenance

16,359+
Curated Q&A Pairs

5,126
Documented Diseases

9
NIH Institutes

132 Words
Avg Clinical Response

Contributing NIH Institutes

- MedlinePlus: Patient education, anatomy, disorders, medications.
- National Cancer Institute (Cancer.gov): Oncology diagnostics & therapies.
- CDC: Communicable diseases, public health guidelines.
- NIDDK: Diabetes, kidney disease, endocrine & metabolic disorders.
- NINDS: Neurological disorders, brain pathologies, dementia.
- GARD: Genetic and rare diseases repository.

Question Taxonomy Distribution

- Symptoms & Signs: Clinical indicators, early warning signs, pain profiles.
- Treatment & Management: Pharmacotherapy, surgical options, care plans.
- Causes & Risk Factors: Genetic etiology, epidemiology, lifestyle habits.
- Diagnosis & Tests: Biopsies, blood screening, radiological imaging.
- Prognosis & Outlook: Life expectancy, recovery timelines, survival rates.

HEALTHCARE AI & NLP CAPSTONE

End-to-End System Architecture

1. Data Ingestion & Schema Normalizer

- Auto-discovers dataset path in environment
- Multi-encoding fallback (UTF-8, Latin-1, CP1252)

2. TF-IDF & Clinical Vector Space Index

- Sublinear term frequency TF-IDF vectorizer
- N-gram feature extraction (unigrams & bigrams)

3. Retrieval & Domain Entity Boosting

- Query vector transformation & cosine distance
- Pathological & anatomical entity detection

4. Clinical Post-Processing & UI

- Clinical Takeaways Extractor (heuristic scoring)
- Plain-English medical glossary matching
- Streamlit responsive reactive web application
- Standalone HTML5 + Tailwind web dashboard

Information Retrieval & NLP Formulation

Vector Space Modeling & TF-IDF

- Sublinear Term Frequency Formulation:
  tf_sub(t, d) = 1 + log(tf(t, d)) if tf > 0, else 0
  Prevents long medical paragraphs from artificially skewing relevance.
- Inverse Document Frequency (IDF):
  idf(t) = log((1 + N) / (1 + df(t))) + 1
  Penalizes ubiquitous medical stopwords (e.g. 'patient', 'disease').
- N-gram Windowing (1, 2):
  Captures compound clinical phrases such as 'blood pressure', 'heart attack', and 'optic nerve'.

Cosine Similarity & Entity Boosting

- Cosine Projection:
  CosineSim(Q, D) = (Q · D) / (||Q|| × ||D||)
  Evaluates normalized directional angle between query and document vectors.
- Clinical Entity Boosting:
  Score(d) = CosineSim(Q, D) + α · EntityMatch(Q, FocusArea_d)
  Empirical boost factor α = 0.08 elevates exact condition matches to top rank.
- Confidence Normalization:
  ConfidencePct = min(100.0, Score(d) × 100)
  Provides transparent match percentages to end users.

Key Features & Clinical Intelligence Modules

Intelligent Natural Language Search

- Accepts free-form patient queries & symptoms
- Instant sub-second response times across 16,000+ records
- Confidence badge ranking & attribution tags

Automated Takeaway Extractor

- Heuristic sentence evaluation based on high-signal clinical verbs
- Synthesizes 3-4 actionable key bullet points
- Saves time for clinicians and patients during consultation

Clinical Jargon Simplifier

- Built-in verified medical terminology lexicon
- Matches and defines terms: Dyspnea, Ischemia, Neuropathy, etc.
- Translates medical terminology into plain English

Disease & Condition Explorer

- Comprehensive catalog across 5,126 conditions
- Categorized accordion view by diagnosis, treatment, prognosis
- Interactive search and filter capabilities

User Experience & Clinical UI Dashboard

Web Application Capabilities

- Streamlit Reactive Frontend: Real-time query execution with zero page reloads, cached dataset loads, and instant state management.
- Prompt Pill Suggestions: One-click inquiry buttons for common clinical scenarios (Asthma, Diabetes, Hypertension, Glaucoma, Heart Attack).
- Dynamic Sidebar Controls: Interactive sliders for confidence thresholds and result counts, with dropdowns for question types and NIH sources.
- Responsive Standalone Webpage: Clean HTML5 + Tailwind CSS dashboard (`index.html`) executable in any browser.

Clinical Design System

- Healthcare Color Palette: Deep teal (#0d9488), navy dark (#0f172a), and clean white surfaces to instill clinical trust.
- Color-Coded Badges: Immediate visual classification for confidence score, medical focus area, question category, and source institute.
- Structured Result Cards: Direct authoritative answer, collapsible takeaway list, glossary tooltips, and related questions.
- Prominent Clinical Disclaimer: HIPAA and educational notices displayed prominently for safety.

Dataset Analytics & Exploratory Data Analysis

Key Quantitative Findings

- Extensive Coverage: Over 16,359 distinct Q&A pairs covering 5,126 conditions, making it one of the largest verified NIH benchmarks.
- Top Documented Conditions: Breast cancer, Type 2 diabetes, Asthma, Glaucoma, Hypertension, and Parkinson's disease.
- Substantial Clinical Depth: Average answer length of 132 words, providing comprehensive guidance rather than superficial snippets.
- Institutional Distribution: MedlinePlus and Cancer.gov constitute the primary knowledge contributors, backed by CDC guidelines.

Analytical Tools & Export Capabilities

- Live Metric Telemetry: Real-time calculation of dataset distributions, total conditions, and word length variance.
- Interactive Charting: Native Altair & Streamlit bar graphs displaying frequency of question types and condition coverage.
- Interactive Data Table: In-app search and filtering of raw clinical pairs across all attributes.
- One-Click CSV Export: Filtered subsets can be downloaded directly from the dashboard for downstream research.

Deployment, CI/CD & Cloud Infrastructure

Cloud Deployment Strategies

- Streamlit Community Cloud: 1-click cloud deployment directly connected to GitHub repository with SSL & global CDN.
- Docker Containerization: Pre-configured Dockerfile (`python:3.11-slim`) with healthcheck probes for scalable Kubernetes/ECS deployment.
- Render Blueprint: Ready-to-deploy `render.yaml` specification for zero-configuration web service deployment.
- GitHub Pages Hosting: Automated GitHub Actions workflow publishing the standalone client web UI (`index.html`).

Code Quality & Version Control

- Git Version Control: Configured with automated push scripts, clean branch workflows, and standardized `.gitignore` rules.
- Automated Test Suite (`test_pipeline.py`): Validates schema normalization, TF-IDF ranking, takeaway extraction, and analytics.
- Portable Tooling: Integrated portable MinGit execution for frictionless deployment on restricted environments.
- Comprehensive Documentation: Full `README.md` containing architectural walkthroughs and deployment guides.

CONCLUSION & FUTURE HORIZONS

Summary of Achievements & Next Steps

Project Achievements

✓ Engineered end-to-end Healthcare QA system over 16,000+ NIH records.

✓ Built high-speed TF-IDF retrieval index with clinical entity boosting.

✓ Implemented automated takeaway summarizer and jargon simplifier.

✓ Designed responsive clinical web UI (Streamlit & standalone HTML).

✓ Packaged with Docker, GitHub CI/CD, and test suite.

Future Roadmap

→ Dense Neural Retrieval: Integrate BioBERT and ClinicalBERT dense vector embeddings for deep semantic matching.

→ Voice Navigation: Add speech-to-text input and audio answer synthesis for accessibility.

→ Multilingual Translation: Real-time localization for non-English speaking patient populations.

→ EHR Integration: Interoperability with Fast Healthcare Interoperability Resources (FHIR) standards.